


Can the model get through the level like a person does?



Agent: best.pt of each level's highest-scoring specialist run; 5 rollouts, stochastic starts (seeds 4242-4246); greedy action selection, no root noise. Human: successful run-throughs from the converted CNeuroMod corpus, same frame-skip. Levels humans never played show no orange band. Not a like-for-like rate comparison — a few rollouts from a hand-picked checkpoint against every recorded human attempt. Generated 2026-08-21 18:46 UTC · git bdeff18